

Instrumenting Python Applications with Prometheus

Emily Woods & Jessica Greene





Emily Woods

@sometimes_milo



- Site Reliability Engineer
- Previously a chemical engineer
- I <3 Python because of its community
- Sewing, plants and books





Jessica Greene

@sleepypioneer



- ★ Software Engineer
 - Self Taught/ Community Taught
- ★ Career Changer
 - previously a coffee roaster
 - & Camera assistant
- ★ Fell in love with Python's versatility
- ★ Community Organiser, Climate activist, knitter, book lover and plant mom



Agenda:

- What is monitoring, why is it important
- Walk through the project we will work on today
- Exposing metrics with prometheus client
Challenge 1: expose base app metrics on /metrics endpoint
- Adding custom metrics, what makes a meaningful metric?
Challenge 2: add a custom metric with labels

BREAK

- Inspecting metrics with Prometheus & Grafana
Challenge 3: Query your metric with PromQL
Challenge 4: Create a Grafana Dashboard

BONUS

- Data types in prometheus, what do we want to monitor?
Take home challenge: make a histogram for request latencies



Monitoring: Why?




Image credit: <https://getnave.com/blog/kanban-metrics/>



Monitoring: What?


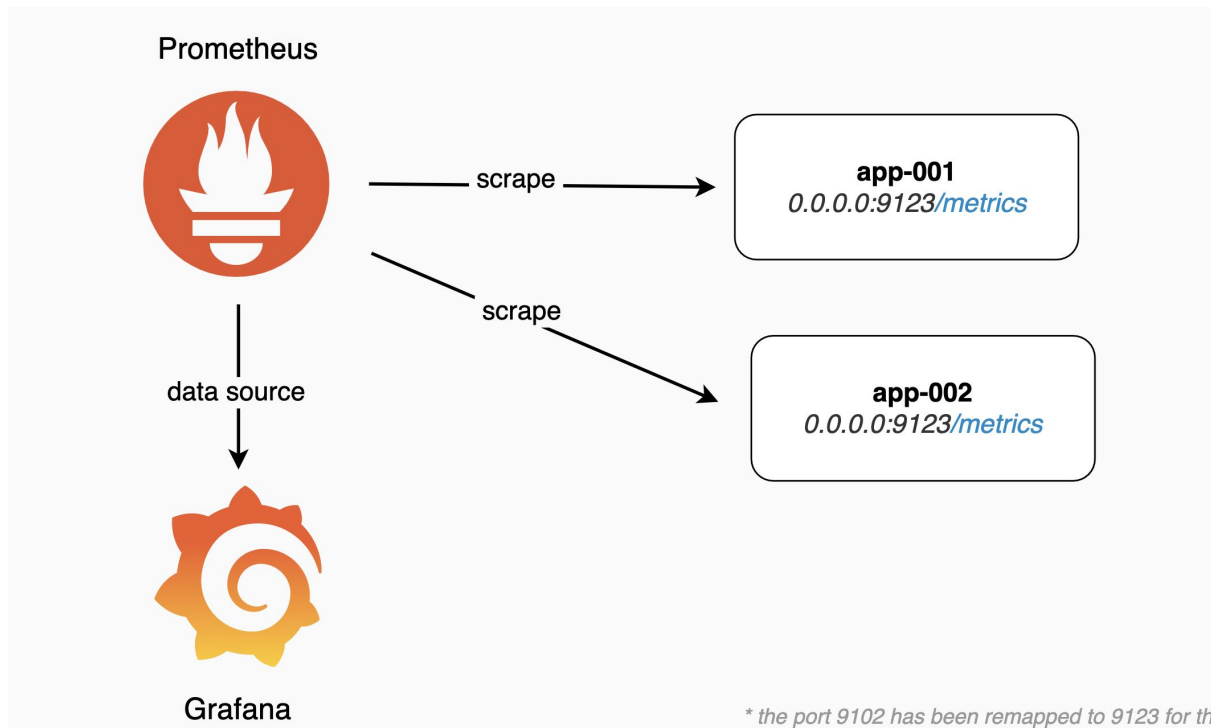


Image credit: <https://www.haproxy.com/blog/haproxy-exposes-a-prometheus-metrics-endpoint/>





Monitoring: How?





Instrumenting your Python Application

<https://github.com/ecosia/python-prometheus-workshop>



What is a metric? (base metrics)

HELP process_cpu_seconds_total **Total user and system CPU time spent in seconds.**

TYPE process_cpu_seconds_total counter

process_cpu_seconds_total 1.01

HELP python_gc_collections_total **Number of times this generation was collected**

TYPE python_gc_collections_total counter

python_gc_collections_total{generation="0"} 53.0

python_gc_collections_total{generation="1"} 4.0

python_gc_collections_total{generation="2"} 0.0

The Python garbage collector has three generations in total, and an object moves into an older generation whenever it survives a garbage collection process on its current generation.

HELP process_start_time_seconds **Start time of the process since unix epoch in seconds.**

TYPE process_start_time_seconds gauge

process_start_time_seconds 1.60251632472e+09



Challenge 1 - expose metrics

1. Import the python client & use the prometheus handler as a base for your request handler
2. run the application a few times and check /metrics



Defining Custom Metrics

- Useful when 'out of the box' metrics aren't sufficient
- To define one, choose a data type and provide:
 - Base name
 - Description
 - Labels

```
7 from prometheus_client import Counter, MetricsHandler
8
9 c = Counter('requests_total', 'requests', ['status', 'endpoint'])
10
11
12 HOST_NAME = '0.0.0.0' # This will map to available port in docker
13 PORT_NUMBER = 8001
14 trees_api_url = "https://api.ecosia.org/v1/trees/count"
15 with open('./templates/treeCounter.html', 'r') as f:
16     html_string = f.read()
17 html_template = Template(html_string)
18
19 def fetch_tree_count():
20     r = requests.get(trees_api_url)
21     # Here is one possible place you may decide to call this metric from
22     c.labels(status=f'{r.status_code}', endpoint='/trees').inc()
23     if r.status_code == 200:
24         return r.json()['count']
25     return 0
26
```



Defining Custom Metrics

- Useful when 'out of the box' metrics aren't sufficient
- To define one, choose a data type and provide:

- Base name
- Description
- Labels

```
7 from prometheus_client import Counter, MetricsHandler
8
9 c = Counter('requests_total', 'requests', ['status', 'endpoint'])
10
11
12 HOST_NAME = '0.0.0.0' # This will map to available port in docker
13 PORT_NUMBER = 8001
14 trees_api_url = "https://api.ecosia.org/v1/trees/count"
15 with open('./templates/treeCounter.html', 'r') as f:
16     html_string = f.read()
17 html_template = Template(html_string)
18
19 def fetch_tree_count():
20     r = requests.get(trees_api_url)
21     # Here is one possible place you may decide to call this metric from
22     c.labels(status=f'{r.status_code}', endpoint='/trees').inc()
23     if r.status_code == 200:
24         return r.json()['count']
25     return 0
26
```



Defining Custom Metrics

- Useful when 'out of the box' metrics aren't sufficient
- To define one, choose a data type and provide:

- Base name
- Description
- Labels

```
7 from prometheus_client import Counter, MetricsHandler
8
9 c = Counter('requests_total', 'requests', ['status', 'endpoint'])
10
11
12 HOST_NAME = '0.0.0.0' # This will map to available port in docker
13 PORT_NUMBER = 8001
14 trees_api_url = "https://api.ecosia.org/v1/trees/count"
15 with open('./templates/treeCounter.html', 'r') as f:
16     html_string = f.read()
17 html_template = Template(html_string)
18
19 def fetch_tree_count():
20     r = requests.get(trees_api_url)
21     # Here is one possible place you may decide to call this metric from
22     c.labels(status=f'{r.status_code}', endpoint='/trees').inc()
23     if r.status_code == 200:
24         return r.json()['count']
25     return 0
26
```



Defining Custom Metrics

- Useful when 'out of the box' metrics aren't sufficient
- To define one, choose a data type and provide:

- Base name
- Description
- Labels

```
7 from prometheus_client import Counter, MetricsHandler
8
9 c = Counter('requests_total', 'requests', ['status', 'endpoint'])
10
11
12 HOST_NAME = '0.0.0.0' # This will map to available port in docker
13 PORT_NUMBER = 8001
14 trees_api_url = "https://api.ecosia.org/v1/trees/count"
15 with open('./templates/treeCounter.html', 'r') as f:
16     html_string = f.read()
17 html_template = Template(html_string)
18
19 def fetch_tree_count():
20     r = requests.get(trees_api_url)
21     # Here is one possible place you may decide to call this metric from
22     c.labels(status=f'{r.status_code}', endpoint='/trees').inc()
23     if r.status_code == 200:
24         return r.json()['count']
25     return 0
26
```



Defining Custom Metrics

- Useful when 'out of the box' metrics aren't sufficient
- To define one, choose a data type and provide:

- Base name
- Description
- Labels

```
7 from prometheus_client import Counter, MetricsHandler
8
9 c = Counter('requests_total', 'requests', ['status', 'endpoint'])
10
11
12 HOST_NAME = '0.0.0.0' # This will map to available port in docker
13 PORT_NUMBER = 8001
14 trees_api_url = "https://api.ecosia.org/v1/trees/count"
15 with open('./templates/treeCounter.html', 'r') as f:
16     html_string = f.read()
17 html_template = Template(html_string)
18
19 def fetch_tree_count():
20     r = requests.get(trees_api_url)
21     # Here is one possible place you may decide to call this metric from
22     c.labels(status=f'{r.status_code}', endpoint='/trees').inc()
23     if r.status_code == 200:
24         return r.json()['count']
25     return 0
26
```



Querying Custom Metrics

```
# HELP requests_total Total requests
```

```
# TYPE requests_total counter
```

```
requests_total{method="get",status="200"} 1.0
```

Description



Querying Custom Metrics

```
# HELP requests_total Total requests
```

```
# TYPE requests_total counter
```

```
requests_total{method="get",status="200"} 1.0
```

Measurement type



Querying Custom Metrics

```
# HELP requests_total Total requests
```

```
# TYPE requests_total counter
```

```
requests_total{method="get",status="200"} 1.0
```

Base name



Custom Metrics

```
# HELP requests_total Total requests
```

```
# TYPE requests_total counter
```

```
requests_total{method="get",status="200"} 1.0
```

Labels



Challenge 2 - custom metric

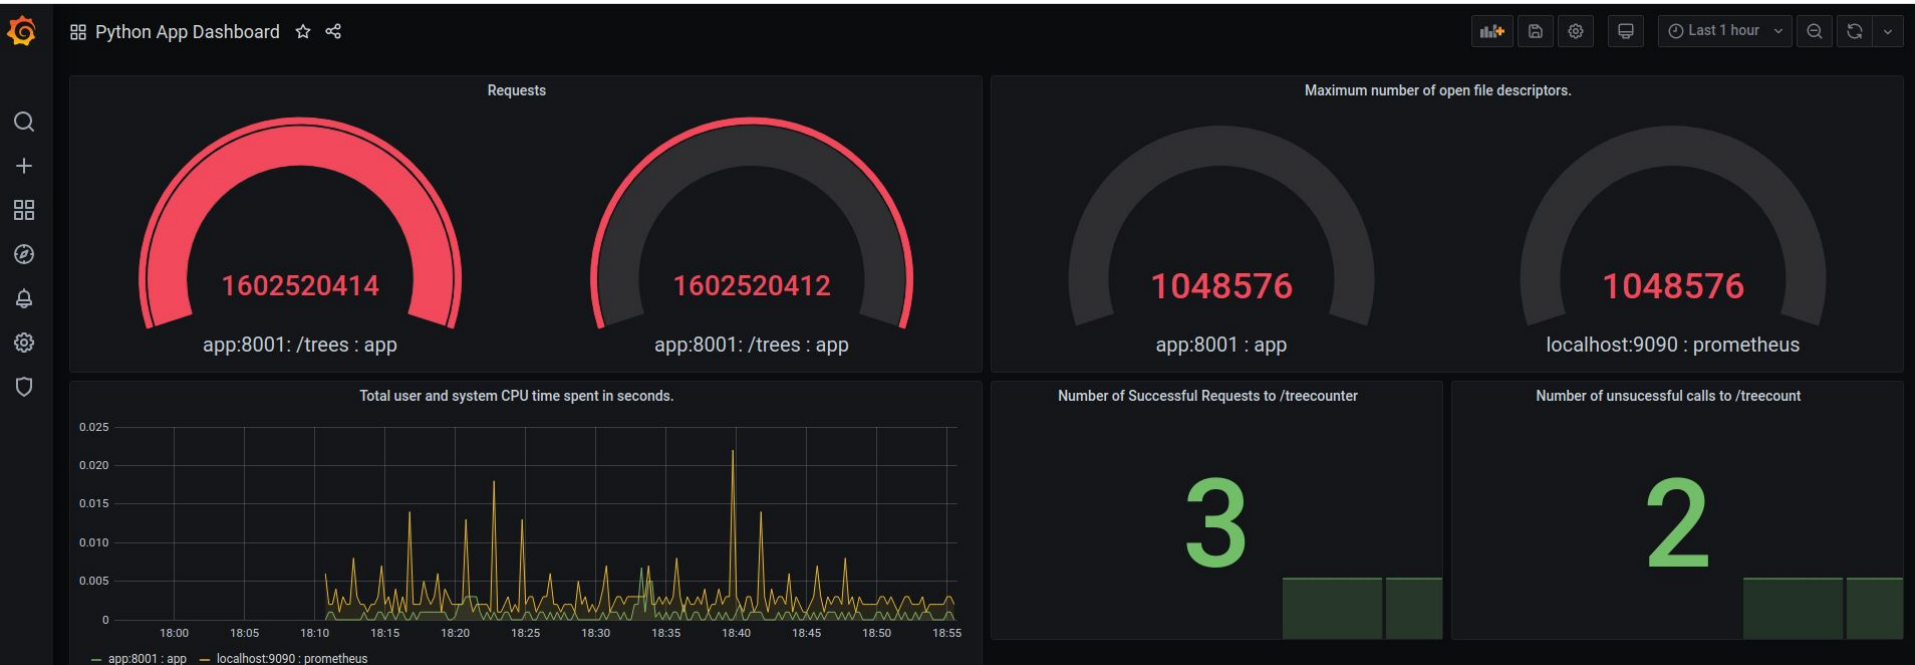
1. add request counter with status code as label & look at /metrics
2. run the application a few times and check again /metrics
3. Add a label to your metric for status code



BREAK



Scraping metrics & creating dashboards





App:

<http://localhost:8001>

Prometheus:

<http://localhost:9090>

Grafana:

<http://localhost:3000>

Monitoring your metrics

<https://github.com/ecosia/python-prometheus-workshop>

docker-compose up



Status > Targets

Prometheus Alerts Graph Status ▾ Help

Targets

All Unhealthy

app (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://app:8001/metrics	UP	instance="app:8001" job="app"	6.785s ago	1.81ms	

prometheus (1/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://localhost:9090/metrics	UP	instance="localhost:9090" job="prometheus"	922ms ago	15.97ms	



Using PromQL to query Metrics

Prometheus Alerts Graph Status ▾ Help

☐ Enable query history [Try experimental React UI](#)

requests_total{endpoint="/trees", status="200"}

Execute - insert metric at cursor - ▾

Graph Console

◀ Moment ▶

Element	Value
requests_total{endpoint="/trees", instance="app:8001", job="app", status="200"}	3

Add Graph

Load time: 35ms
Resolution: 14s
Total time series: 1

[Remove Graph](#)

requests_total{endpoint="/trees", status="200"}



Creating a panel in your dashboard

The screenshot shows the Grafana 'New dashboard / Edit Panel' interface. The main panel area displays the title 'Number of Successful Requests to /treecounter' and a large green number '3'. A red box highlights the visualization type 'Stat' in the right-hand sidebar, which also shows a preview of the 'Stat' visualization with the value '12.4'. The bottom of the interface shows the query editor with the query `requests_total{endpoint="/trees", status="200"}` and the 'Query options' section. The top right corner features buttons for 'Discard', 'Save', and 'Apply'.

← New dashboard / Edit Panel

Fill Fit Exact Last 6 hours

Number of Successful Requests to /treecounter

3

Query 1 Transform 0

default Query options MD = auto = 1470 Interval = 15s Query inspector

requests_total{endpoint="/trees", status="200"}

Legend Min step Resolution 1/1 Format Time series Instant Prometheus

+ Query

Panel Field Overrides

Settings

Panel title
Number of Successful Requests to /treecounter

Description
Panel description supports markdown and links.

Transparent
Display panel without a background.

Visualization

Filter visualizations

Graph Stat 12.4

Gauge Bar gauge

Table Text

Challenge 3 - Query your metric with PromQL *(section 3 of readme)*

1. Run the app and prometheus with docker-compose
2. Query your metrics in the prometheus UI using PromQL



Challenge 4 - Create a Grafana Dashboard

(section 3 of readme)

1. Run the app, prometheus and Grafana with docker-compose
2. Go to localhost:3000
3. Create a Grafana Dashboard
4. Create a panel to visualise your customer metric (split by labels, irate or cumulative)





Data types in Prometheus




Created by Vectors Point
from Noun Project

1. Counter



Created by IconMark
from Noun Project

2. Gauge



Created by Nhor
from Noun Project

3. Histogram


4. Summary





Histogram

- In Prometheus, a histogram measures the frequency of value observations that fall into **buckets**.
- Useful when you want approximations over a known range of values
- Common uses: measuring response duration, request size





What a Histogram exposes

- **Cumulative counters** for the observation buckets, exposed as `<metric_name>_bucket`

```
request_latency_seconds_bucket{endpoint="/treecounter",le="0.1"} 1.0
```

```
request_latency_seconds_bucket{endpoint="/treecounter",le="0.5"} 4.0
```

```
request_latency_seconds_bucket{endpoint="/treecounter",le="1.0"} 5.0
```

```
request_latency_seconds_bucket{endpoint="/treecounter",le="+Inf"} 5.0
```

- The **total sum** of all observed values, exposed as `<metric_name>_sum`

```
request_latency_seconds_sum{endpoint="/treecounter"} 1.14
```

- The **count** of events that have been observed, exposed as `<metric_name>_count` (identical to `<metric_name>_bucket{le="+Inf"}`)

```
request_latency_seconds_count{endpoint="/treecounter"} 5.0
```

Bonus Challenge

1. make a histogram for request latencies.

*(hint for measuring latencies
with time)*



Solution in branch: *solution_challenge_5*



@ecosia
@sometimes_milo
@sleepypioneer

Go forth and monitor!



Gif credit: <https://giphy.com/tonybabel>





Resources

<https://prometheus.io>

<https://prometheus.io/docs/practices/histograms/>

<https://grafana.com/>

<https://tomgregory.com/the-four-types-of-prometheus-metrics/>

<https://github.com/ecosia/python-prometheus-workshop>

